

Introduction to Parallelism & Parallelism on HPC

Examples in Julia and Python

Criston Hyett

December 7, 2022

What is parallelism?

- ▶ The common idea of **divide and conquer**
- ▶ Works extremely well in many relevant scenarios
 - ▶ Large dimension linear algebra, (i.e. ML)
 - ▶ Monte Carlo
 - ▶ Ordinary/Partial/Stochastic Differential equations (big linear algebra + Monte Carlo)
- ▶ Not a panacea!
 - ▶ Inefficient code can be inefficient *and* consume plenty of cpu-hours on the HPC

When should you care?

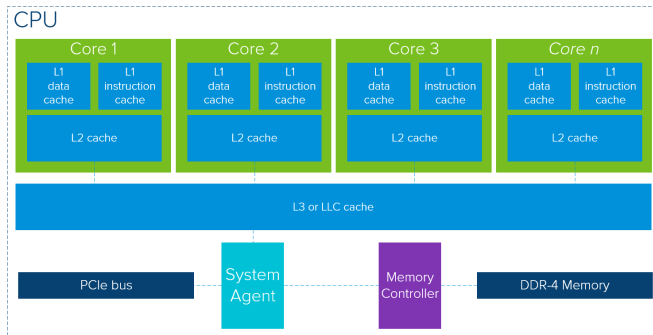
- ▶ When the benefits outweigh the costs!
- ▶ Costs
 - ▶ **Human costs:** Coding, debugging, refactoring
 - ▶ **Computational Costs:** parallelism requires coordination between threads, and/or nodes. When this coordination is required often, or significant data is passing between threads, parallel benefits may be outweighed.
- ▶ Benefits
 - ▶ **Computation speed**
 - ▶ **Smaller memory overheads per thread**
 - ▶ Process independence

When should you care?


- ▶ When the benefits outweigh the costs!
- ▶ Costs
 - ▶ **Human costs:** Coding, debugging, refactoring
 - ▶ **Computational Costs:** parallelism requires coordination between threads, and/or nodes. When this coordination is required often, or significant data is passing between threads, parallel benefits may be outweighed.
- ▶ Benefits
 - ▶ **Computation speed**
 - ▶ **Smaller memory overheads per thread**
 - ▶ Process independence

Many of the costs are mitigated by modern languages providing thread-safe, optimized libraries.

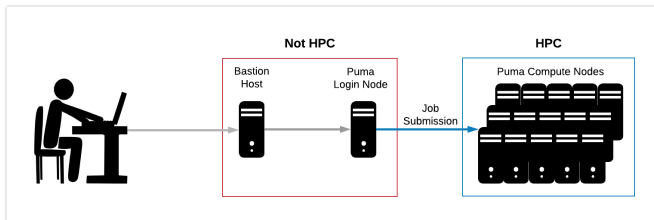
Parallel paradigms: CPU



Parallel paradigms: GPU



Parallel paradigms: Multi-node



<https://public.confluence.arizona.edu/display/UAHPC/Puma+Quick+Start>

Word of warning

- ▶ Not all code is parallelizable! Improper implementations can lead to
 - ▶ Race conditions
 - ▶ Deadlock
 - ▶ Memory Corruption

```
badParallel.jl
```

```
a = 0;
for i in 1:1000
    global a += 1;
end
println("serial a = $(a)");

a = 0;
Threads.@threads for i in 1:1000
    global a += 1;
end
println("parallel a = $(a)");

# serial a = 1000
# parallel a = 881
```


CPU Thread parallelism

Julia: helloWorld.jl

```
numThreads = Threads.nthreads();
Threads.@threads for i in 1:numThreads
    println("Hello World!"*
        "This is thread # $(Threads.threadid())");
end

# Hello World! This is thread # 1
# Hello World! This is thread # 6
# Hello World! This is thread # 3
# Hello World! This is thread # 4
# Hello World! This is thread # 5
# Hello World! This is thread # 2
```

Python: helloWorld.py

```
import threading;
from time import sleep;
import numpy as np;

def helloWorld():
    sleep(np.random.random());
    print("Hello world! This is {}".format(threading.current_thread().name))

if __name__ == "__main__":
    for i in range(6):
        t = threading.Thread(
            target=helloWorld,args=[]);
        t.start()

# Hello world! This is Thread-4
# Hello world! This is Thread-1
# Hello world! This is Thread-6
# Hello world! This is Thread-2
# Hello world! This is Thread-5
# Hello world! This is Thread-3
```

GPU Thread parallelism

- ▶ Even easier introduction to Cuda

Node parallelism via Slurm

monteCarlo.jl

```
basepath = "/home/u1/cmhyett/.julia/dev/LDM/"
using Pkg;
Pkg.activate(basepath);
Pkg.instantiate();
denom = ARGS[1];
outputPath = ARGS[2];
include(basepath * "./src/LDM.jl");
inputPath = basepath * "./results/pi$(denom)/lagrData_pi$(denom).jls"
LDM.TBNN.runTBNN(inputPath,
                  outputPath,
                  maxiters=800,
                  learningRate=1e-3)
```

Node parallelism via Slurm

monteCarlo.slurm

```
#!/bin/bash
```

```
# -----
```

```
### PART 1: Requests resources to run your job.
```

```
# -----
```

```
### Optional. Set the job name
```

```
#SBATCH --job-name=test_test
```

```
### Optional. Set the output filename.
```

```
### SLURM reads %x as the job name and %j as the job ID
```

```
#SBATCH --output=%x-%j.out
```

```
### REQUIRED. Specify the PI group for this job
```

```
###SBATCH --account=chertkov
```

```
### Optional. Request email when job begins and ends
```

```
### SBATCH --mail-type=ALL
```

```
### Optional. Specify email address to use for notification
```

```
### SBATCH --mail-user=<YOUR NETID>@email.arizona.edu
```

```
### REQUIRED. Set the partition for your job.
```

```
#SBATCH --partition=windfall
```

```
### REQUIRED. Set the number of cores that will be used for this job.
```

```
#SBATCH --ntasks=20
```

```
### REQUIRED. Set the number of nodes
```

```
#SBATCH --nodes=1
```

```
### REQUIRED. Set the memory required for this job.
```

```
#SBATCH --mem=64gb
```

```
### REQUIRED. Specify the time required for this job, hhh:mm:ss
```

```
#SBATCH --time=02:00:00
```

```
#SBATCH --array=1-20
```

```
# -----
```

```
### PART 2: Executes bash commands to run your job
```

```
# -----
```

```
### Load required modules/libraries if needed
```

```
module load python/3.8
```

```
module load julia/1.7.2
```

```
### change to your script's directory
```

```
cd /home/u1/cmhyett/.julia/dev/LDM/results/pi512/
```

```
julia ../tbnnFamilyTrainScript.jl 512 /xdisk/chertkov/cmhyett/tbnnFamilyTraining/pi512/latent_space/task_$(SLURM_ARRAY_TASK_ID)/
```

References & Further Reading

- ▶ <https://core.vmware.com/resource/exploring-gpu-architecture>
- ▶ <https://public.confluence.arizona.edu/display/UAHPC/Puma+Quick+Start>
- ▶ <https://www.tensorflow.org/guide/gpu>
- ▶ <https://developer.nvidia.com/blog/even-easier-introduction-cuda/>
- ▶ <https://public.confluence.arizona.edu/display/UAHPC/Using+and+Installing+Python>